

Lecture 9

Transaction Management

Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

Introduction

- **Collections of operations** that form a **single logical unit of database processing** are called **Transactions**.
- A **transaction** includes one or more **database access operations**—these can include **insertion, deletion, modification, or retrieval operations**.
- If the **database operations** in a transaction **do not update** the **database** but **only retrieve data**, the **transaction is called a read-only transaction**; otherwise it is known as a **read-write transaction**.
- To **ensure integrity** of the data, we require that the **database system** maintain the following **properties of the transactions**:
 - **Atomicity**
 - **Consistency**
 - **Isolation**
 - **Durability**

Properties of Transaction

- **ATOMICITY**

- Either **all operations** of the **transaction** are **reflected properly** in the **database**, or **none are**.

- **CONSISTENCY**

- Execution of a **transaction** in **isolation** (that is, with no other transaction executing concurrently) **preserves the consistency** of the **database**.

- **ISOLATION**

- Even though **multiple transactions** may **execute concurrently**, the system guarantees that, for every **pair of transactions** T_i and T_j ,
 - It appears to T_i that either T_j **finished execution** before T_i **started**, or
 - T_j **started execution** after T_i **finished**.
 - Thus, **each transaction** is unaware of other **transactions executing concurrently** in the system.

Properties of Transaction

- **DURABILITY**
- After a **transaction completes successfully**, the **changes it has made** to the **database persist**, even if there are **system failures**.
- **Note:** These **properties of transactions** are also known as **ACID properties** of transaction.

Properties of Transaction: Example

- Consider a simplified **banking system** consisting of **several accounts** and a set of **transactions** that **access** and **update** those **accounts**.
- We assume that the **database permanently resides on disk**, but that **some portion** of it is temporarily residing in **main memory buffer**.
- **A Transactions** can access **data** using **two operations**:
 - **read(X)**, which **transfers the data item X** from the **database** to a **local buffer** belonging to the transaction that executed the **read operation**.
 - **write(X)**, which **transfers the data item X** from the **local buffer** of the transaction that executed the **write back to the database**.

Properties of Transaction: Example

- Let T_i be a Transaction that transfers \$50 from account A to account B .
- This transaction can be defined as:
 - T_i : read(A);
 - $A := A - 50$;
 - write(A);
 - read(B);
 - $B := B + 50$;
 - write(B).

Properties of Transaction: Example

- **ATOMICITY**
- Suppose that, just **before the execution** of **transaction T_i** the values of **accounts A** and **B** are **\$1000** and **\$2000**, respectively.
- Now suppose that, during the execution of **transaction T_i** , a **failure occurs** that prevents **T_i** from completing its execution successfully.
- Examples of such failures include **power failures, hardware failures, and software errors.**
- Further, suppose that the **failure happened after the $\text{write}(A)$ operation** but **before the $\text{write}(B)$ operation.**
- In this case, the values of **accounts A** and **B** reflected in the **database** are **\$950** and **\$2000**.
- The system **destroyed \$50** as a result of this **failure.**

Properties of Transaction: Example

- In particular, we note that the sum $A + B$ is no longer preserved.
- The basic idea behind **ensuring atomicity** is this: The **database system keeps track (on disk) of the old values** of any **data** on which a transaction performs a **write**.
- And, if the **transaction does not complete its execution**, the **database system restores the old values** to make it **appear as though the transaction never executed**.
- Ensuring **Atomicity** is handled by a component called the **Transaction-management component** of **DBMS** software .

Properties of Transaction: Example

- **CONSISTENCY**
- The **Consistency** requirement here is that the **sum of A and B** be **unchanged** by the **execution** of the **transaction**.
- Without the **consistency requirement**, **money** could be **created or destroyed** by the **transaction**!
- Ensuring **consistency** for an individual transaction is the **responsibility** of the **Application programmer** who **codes the transaction**.

Properties of Transaction: Example

- **ISOLATION**
- If **several transactions** are **executed concurrently**, their operations may **interleave** in some **undesirable way**, resulting in an **inconsistent state**.
- For **example** the **database** is temporarily **inconsistent** while the transaction to transfer **funds from A to B** is executing, with the **deducted total written to A** and **the increased total yet to be written to B**.
- If a **second concurrently running transaction** reads **A** and **B** at this **intermediate point** and computes **$A+B$** , it will observe an **inconsistent value**.
- Furthermore, if this **second transaction** then performs updates on **A** and **B** based on the **inconsistent values** that it read, the **database** may be left in an **inconsistent state** even after both transactions have completed.

Properties of Transaction: Example

| T1 | T2 |
|--|--|
| read(A) A=A-50 write(A) Read(B) B=B+50 Write(B) | |
| | read(A) temp= A*0.1 A= A – temp write(A) Read(B) B=B+temp write(B) |



| T1 | T2 |
|---|--|
| read(A) A=A-50 | |
| | read(A) temp= A*0.1 A= A – temp write(A) Read(B) |
| write(A) Read(B) B=B+50 write(B) | |
| | B=B+temp write(B) |

Properties of Transaction: Example

- A way to avoid the problem of concurrently executing transactions is to execute **transactions serially**—that is, one after the other.
- However, **concurrent execution of transactions provides significant performance benefits.**
- The **Isolation property** of a **transaction ensures** that the **concurrent execution of transactions results** in a **system state** that is **equivalent** to a **state** that could have been obtained had these **transactions executed one at a time in some order.**
- Ensuring the **isolation property** is the **responsibility** of a **component** of the **database system** called the **Concurrency-control component.**

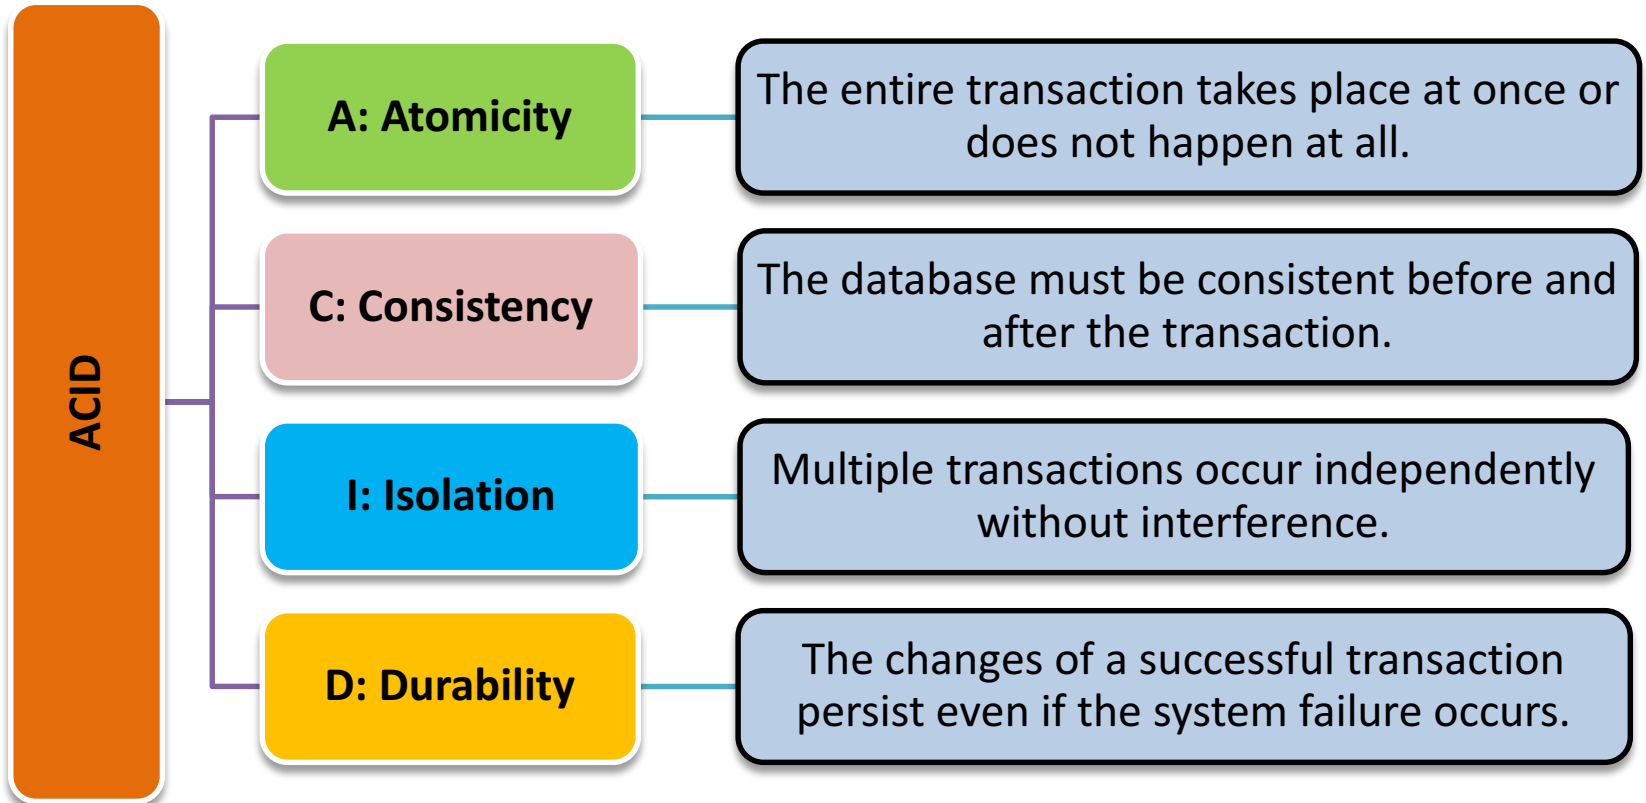
Properties of Transaction: Example

- **DURABILITY**
- Once the **execution** of the **transaction completes successfully**, and the user who initiated the transaction has been notified that the transfer of funds has taken place.
- It must be the case that **no system failure will result in a loss of data corresponding to this transfer of funds.**
- The **durability property** guarantees that, **once a transaction completes successfully**, all the **updates** that it carried out on the **database persist**, even if there is a **system failure** after the **transaction completes** execution.

Properties of Transaction: Example

- We assume that a **failure of the computer system** may result in **loss of data in main memory**, but **data** written to **disk** are **never lost**.
- We can guarantee **Durability** by ensuring that either:
 - The **updates carried out by the transaction** have been **written to disk** before the **transaction completes**.
 - **Information about** the **updates** carried out by the **transaction** and written to **disk** is sufficient to enable the **database** to **reconstruct the updates** when the **database system** is **restarted after the failure**.
- Ensuring **Durability** is the **responsibility** of a **component** of the database system called the **Recovery-management component**.

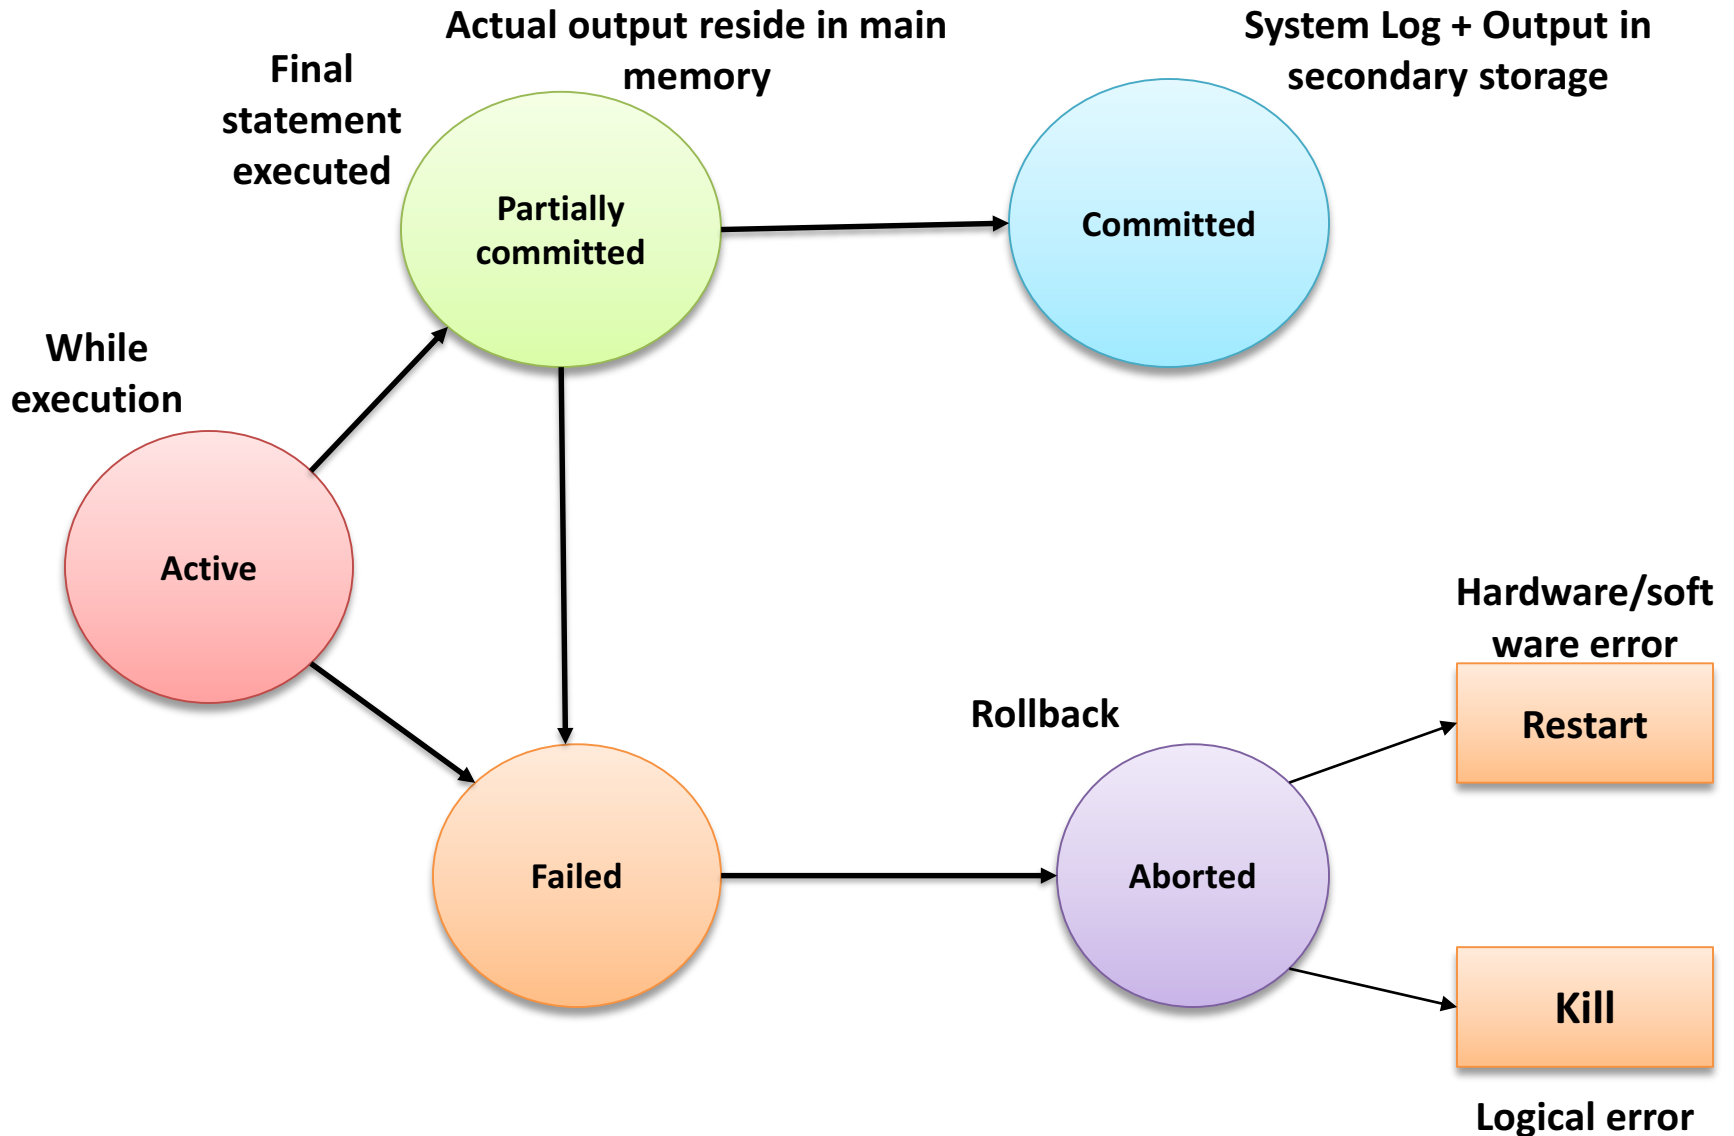
Properties of Transaction: Summary



Transaction States

- A **Transaction** must be in one of the following **states**:
- **Active**
 - The **initial state**; the **transaction** stays in this **state** while it is **executing**.
- **Partially committed**
 - After the **final statement** has been **executed**.
- **Failed**
 - After the discovery that normal execution can no longer proceed.
- **Aborted**
 - After the transaction has been **rolled back** and the database has been **restored** to its **state** prior to the start of the transaction.
- **Committed**
 - After **successful completion**.

Transaction States



Transaction States

- A Transaction starts in the **Active state**.
- When it **finishes its final statement**, it enters the **Partially committed state**.
 - At this point, the **transaction** has **completed** its **execution**,
 - But it is **still possible** that it may have to be **aborted**, since the **actual output** may still be **temporarily residing** in **main memory**, and
 - Thus a **hardware failure** may **preclude** its **successful completion**.
- The **database system** then **writes out** enough information to **disk(log file)** that, even in the event of a **failure**, the **updates** performed by the **transaction** can be **re-created** when the **system restarts** after the **failure**.
- When **the last of this information is written out**, the **transaction** enters the **Committed state**.

Transaction States

- A **Transaction** enters the **Failed state** after the system determines that the **transaction** can **no longer proceed** with its **normal execution** (for example, because of **hardware** or **logical errors**).
- Such a **transaction** must be **Rolled back**, then, it enters the **Aborted state**. At this point, the system has **two options**:
 - **Restart the transaction**
 - If the **transaction** was **aborted** as a result of **some hardware or software error**.
 - A **restarted transaction** is considered to be a **new transaction**.
 - **Kill the transaction**
 - It usually does so because of some **internal logical error** that can be corrected only by **rewriting** the **application program**, or because the **input was bad**, or because the **desired data** were **not found** in the **database**.

The System Log

- To be able to **recover from failures** that affect **transactions**, the system maintains a **log** to **keep track of all transaction operations** that affect the **values** of **database items**.
- The **log** is a **sequential, append-only file** that is **kept on disk**, so it is **not affected** by any type of **failure** except for **disk** or **catastrophic failure**.
- Typically, one (or more) **main memory buffers** hold the **last part of the log file**, so that **log entries** are **first added** to the **main memory** buffer.
- When the **log buffer** is filled, or when certain other conditions occur, the **log buffer** is *appended to the end of the log file on disk*.
- In addition, the **log file** from **disk** is **periodically backed up** to **archival storage (tape)** to guard against **catastrophic failures**.
- The following are the **types of entries**—called **log records**—that are **written** to the **log** file and the corresponding action for each log record:

The System Log

1. [start_transaction, T].

Indicates that transaction T has started execution.

2. [write_item, T , X , *old_value*, *new_value*].

Indicates that transaction T has **changed** the value of database item X from *old_value* to *new_value*.

3. [read_item, T , X].

Indicates that transaction T has **read** the value of database item X .

4. [commit, T].

Indicates that transaction T has **completed successfully**, and affirms that its effect can be committed (recorded permanently) to the database.

5. [abort, T].

Indicates that transaction T has been **aborted**.

Commit Point of a Transaction

- A transaction T reaches its **Commit point** when:
 - All its operations that access the database have been executed successfully *and*
 - the effect of all the transaction operations on the database have been recorded in the log file.
- Beyond the **Commit point**, the **Transaction** is said to be **Committed**, and its effect must be *permanently recorded in the database*.
- The transaction then writes a commit record [**commit**, T] into the log.

Commit Point of a Transaction

- If a system failure occurs:
 - We can search back in the log for all transactions T that have written a **[start_transaction, T]** record into the log but have not written their **[commit, T]** record yet;
 - These transactions may have to be *rolled back to undo their effect* on the database during the recovery process.
 - The Transactions that have written their Commit record **[commit, T]** in the log must also have recorded all their **WRITE** operations in the log, so their effect on the database can be **redone** from the log records.

Serial Execution

- The execution sequences of transactions are called **schedules**, shown below:

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) |

Schedule 1

A serial schedule in which T_2 is followed by T_1 .

- The final values of accounts A and B , after the execution takes place, are \$855 and \$2145, respectively.
- Thus, the **total amount of money in accounts A and B** —that is, the **sum $A + B$** is **preserved** after the execution of both **transactions**.

Serial Execution

- A serial schedule in which T1 is followed by T2.

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) |

Schedule 2

A serial schedule in which T1 is followed by T2.

- Check the **sum $A + B$** is **preserved** or not?
- These schedules(Schedule 1 and Schedule 2) are **serial**.
- Each **serial schedule** consists of a sequence of instructions from various transactions, where the instructions belonging to one single transaction appear together in that schedule.
- Thus, for a **set of n transactions**, there exist **$n!$ different valid serial schedules**.

Concurrent Executions

- **Transactions** submitted by the various users may **execute concurrently** and may **access** and **update** the **same database items**.
- If this **concurrent execution** is *uncontrolled*, it may lead to **problems**, such as an **inconsistent database**.
- **Example:** Consider two **transactions T1** and **T2**:

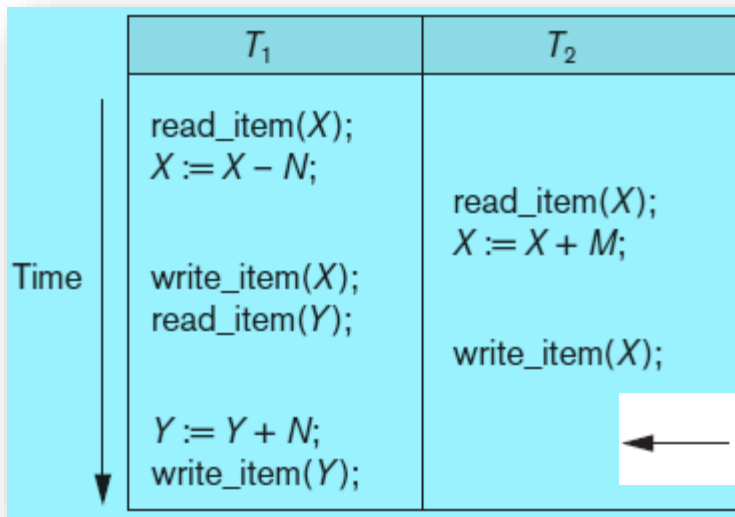
| T_1 |
|--|
| read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y); |

| T_2 |
|---|
| read_item(X); $X := X + M$; write_item(X); |

Problems with Concurrent Executions

- **The Lost Update Problem**

- This **problem** occurs when **two transactions** that access the **same database items** have their **operations interleaved** in a way that makes the value of some database items **incorrect**.
- Suppose that **transactions T_1 and T_2** are **submitted** at approximately the **same time**, and suppose that their operations are **interleaved** as shown in Figure:



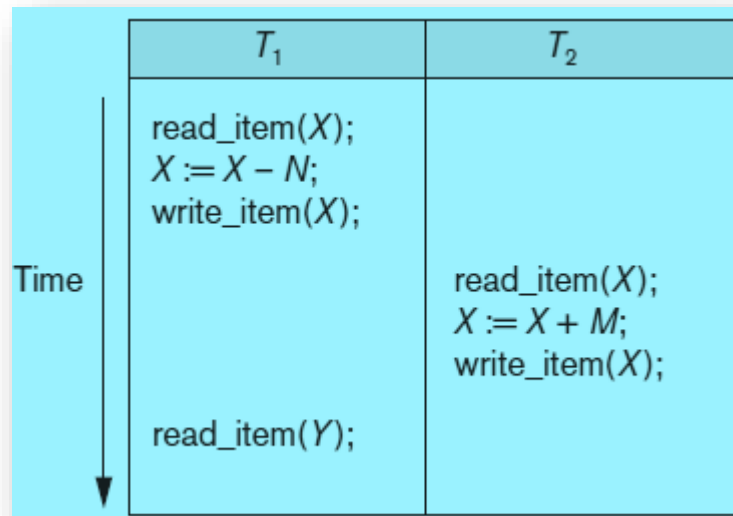
- if $X = 80$ at the start, $N = 5$ and $M = 4$ the final result should be $X = 79$.
- However, in the interleaving of operations it is $X = 84$.
- Because the update in T_1 that deducted five from X was **lost**.

← Item X has an incorrect value because its update by T_1 is **lost** (overwritten).

Temporary Update/Dirty Read Problem

- This **problem occurs** when one **transaction** updates a **database** item and then the **transaction fails** for some reason.
- Meanwhile, the **updated item is accessed (read) by another transaction** before it is changed back to its **original value**.
- Figure shows an **example** where **T1 updates item X** and then **fails before completion**, so the **system** must change **X** back to its **original value**.
- Before it can do so, however, **transaction T2 reads the temporary value of X**, which will not be recorded permanently in the database because of the failure of **T1**.
- The value of **item X** that is read by **T2** is called **dirty data** because it has been created by a transaction that has not completed and **committed** yet; hence, this problem is also known as the **dirty read problem**.

Temporary Update/Dirty Read Problem



Transaction T_1 fails and must change the value of X back to its old value; meanwhile T_2 has read the *temporary* incorrect value of X .

The Incorrect Summary Problem

- If one **transaction** is calculating an **aggregate summary** function on a number of **database items**.
- While other **transactions** are updating some of these items, the **aggregate function** may calculate some values before they are updated and others after they are updated.

| T_1 | T_3 |
|---|--|
| <pre>read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y);</pre> | <pre>sum := 0; read_item(A); sum := sum + A; ⋮ read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre> |

← T_3 reads X after N is subtracted and reads Y before N is added; a wrong summary is the result (off by N).

Concurrent Execution

- If **two transactions** are running **concurrently**, the **operating system** may execute **one transaction** for a little while, then perform a **context switch**, execute the **second transaction** for some time, and then switch back to the first transaction for some time, and so on.
- With **multiple transactions**, the **CPU time is shared** among all the **transactions**.
- Several **execution sequences are possible**, since the various instructions from both transactions may now be **interleaved**.
- Thus, the **number of possible schedules** for a set of **n transactions** is **much larger than $n!$** .
- **Example:** Let **two transactions** are executed **concurrently**.
- One possible schedule is **Schedule 3** shown in figure:

Concurrent Execution

| T ₁ | T ₂ |
|--|--|
| read(<i>A</i>) <i>A</i> := <i>A</i> - 50 write(<i>A</i>) | read(<i>A</i>) <i>temp</i> := <i>A</i> * 0.1 <i>A</i> := <i>A</i> - <i>temp</i> write(<i>A</i>) |
| read(<i>B</i>) <i>B</i> := <i>B</i> + 50 write(<i>B</i>) | read(<i>B</i>) <i>B</i> := <i>B</i> + <i>temp</i> write(<i>B</i>) |

Schedule 3—a concurrent schedule equivalent to **Schedule 1**.

- After this execution takes place, we arrive at the same state as the one in which the transactions are executed serially in the **order T₁ followed by T₂**.
- The **sum *A* + *B*** is indeed **preserved**.

Concurrent Execution

- **Not all concurrent executions result in a correct state.**
- Consider the schedule of Figure given below:

| T_1 | T_2 |
|--|--|
| read(A) $A := A - 50$ | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) |
| write(A) read(B) $B := B + 50$ write(B) | $B := B + temp$ write(B) |

Schedule 4—a concurrent schedule.

- After the execution of this schedule, we arrive at a state where the final values of accounts **A** and **B** are **\$950** and **\$2100**, respectively.
- This final state is an ***inconsistent state***, since we have **gained \$50** in the process of the **concurrent execution**.
- Indeed, the **sum $A + B$** is **not preserved** by the execution of the two transactions.

Concurrent Execution

- We can ensure **consistency** of the **database** under **concurrent execution** by making sure that **any concurrent schedule** that executed has the **same effect** as a **schedule** that could have **occurred without any concurrent execution**.
- That is, *the **concurrent schedule** should, in some sense, be **equivalent to a serial schedule**.*
- **Serializable schedule**
- A **schedule S** of **n** transactions is **serializable** if it is **equivalent** to some **serial schedule** of the **same n** transactions.
- There are **two types of serializability**:
 1. **Conflict Serializability**
 2. **View Serializability**


Conflict Serializability

- Let us consider a **schedule S** in which there are **two consecutive** instructions ***li*** and ***lj***, of transactions ***Ti*** and ***Tj***, respectively.

| Schedule S | |
|------------|----|
| Ti | Tj |
| li | lj |

- If ***li*** and ***lj*** refer to **different data items**, then we can **swap *li* and *lj*** without affecting the results of any instruction in the schedule. **Example:**


| Schedule S | |
|------------|----------|
| Ti | Tj |
| Read(Q) | Write(R) |



| Schedule S | |
|------------|----------|
| Ti | Tj |
| Read(Q) | Write(R) |

Conflict Serializability

- However, if li and lj refer to the same data item Q , then the **order** of the two steps may **matter**.
- There are **four cases** that we need to consider:
- **Case 1: $li = \text{read}(Q)$, $lj = \text{read}(Q)$.**
 - The order of li and lj does **not matter**.
 - Since the same value of Q is read by Ti and Tj , regardless of the order.



Conflict Serializability

- **Case 2: $li = \text{read}(Q)$, $lj = \text{write}(Q)$.**

- If **li** comes before **lj** , then **Ti** does not read the value of **Q** that is written by **Tj** in instruction **lj** .

| Schedule S | |
|------------|----------|
| T_i | T_j |
| Read(Q) | Write(Q) |

| Schedule S | |
|------------|----------|
| T_i | T_j |
| Read(Q) | Write(Q) |

- If **lj** comes before **li** , then **Ti** reads the value of **Q** that is written by **Tj** .
- Thus, the **order of li and lj** matters.


- **Case 3: $li = \text{write}(Q)$, $lj = \text{read}(Q)$.**

- The order of **li** and **lj** matters for reasons similar to those of the **previous case**.

Conflict Serializability

- Case 4: $li = \text{write}(Q)$, $lj = \text{write}(Q)$.

- Since both instructions are **write** operations, the **order** of these instructions **does not affect** either Ti or Tj .



- However, the value obtained by the next **read(Q)** instruction of S is affected, since the result of only the **latter** of the two **write instructions** is **preserved** in the database.
- If there is no other **write(Q)** instruction after li and lj in S , then the order of li and lj directly affects the final value of Q in the database state that results from schedule S .

Conflict Serializability

- Thus, only in the case where both I_i and I_j are read instructions does the relative order of their execution not matter.
- We say that I_i and I_j conflict if they are operations by different transactions on the same data item, and at least one of these instructions is a write operation.

Conflict Serializability

- Consider **Schedule 3**:
- The **write(A)** instruction of **T1** **conflicts** with the **read(A)** instruction of **T2**.
- However, the **write(A)** instruction of **T2** does **not conflict** with the **read(B)** instruction of **T1**, because the two instructions access different data items.
- Let **li** and **lj** be **consecutive instructions** of a **schedule S**.
- If **li** and **lj** are instructions of **different transactions** and **li** and **lj** **do not conflict**, then **we can swap the order of li and lj** to produce a new **schedule S'**.
- Since the **write(A)** instruction of **T2** in **Schedule 3** does not conflict with the **read(B)** instruction of **T1**, we can **swap** these instructions to generate an equivalent schedule, **schedule 5**.

| T ₁ | T ₂ |
|--------------------------------------|---|
| read(A) $A := A - 50$ write(A) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) |
| read(B) $B := B + 50$ write(B) | |
| | read(B) $B := B + temp$ write(B) |

| T ₁ | T ₂ |
|---------------------|---------------------|
| read(A) write(A) | read(A) |
| read(B) | write(A) |
| write(B) | read(B) write(B) |

Conflict Serializability

- We **continue** to **swap non-conflicting instructions**:
 - Swap the `read(B)` instruction of T_1 with the `read(A)` instruction of T_2 .
 - Swap the `write(B)` instruction of T_1 with the `write(A)` instruction of T_2 .
 - Swap the `write(B)` instruction of T_1 with the `read(A)` instruction of T_2 .
- The final result of these swaps, **schedule 6**.

| T_1 | T_2 |
|----------|----------|
| read(A) | |
| write(A) | |
| | read(A) |
| read(B) | |
| | write(A) |
| write(B) | |
| | read(B) |
| | write(B) |

| T_1 | T_2 |
|----------|----------|
| read(A) | |
| write(A) | |
| read(B) | |
| write(B) | |
| | read(A) |
| | write(A) |
| | read(B) |
| | write(B) |

- Thus, we have shown that **schedule 3** is **equivalent** to a **serial schedule**.
- This **equivalence** implies that, regardless of the initial system state, **schedule 3** will produce the same **final state** as will some **serial schedule**.

Conflict Serializability

- If a **schedule S** can be **transformed** into a **schedule S'** by a **series of swaps of non-conflicting instructions**, we say that **S and S'** are **conflict equivalent**.
- In our previous examples, **Schedule 1** is **not conflict equivalent** to **Schedule 2**.
- However, **Schedule 1** is **conflict equivalent** to **Schedule 3**.

| T ₁ | T ₂ |
|--|---|
| read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) |

Schedule 1

| T ₁ | T ₂ |
|--|---|
| read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) |

Schedule 2

| T ₁ | T ₂ |
|--|---|
| read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) |

Schedule 3

Conflict Serializability

- The concept of **conflict equivalence** leads to the concept of **conflict serializability**.
- We say that a concurrent ***schedule S is conflict serializable if it is conflict equivalent to a serial schedule.***
- Thus, **Schedule 3** is **conflict serializable**, since it is **conflict equivalent** to the **serial Schedule 1**.

| T ₁ | T ₂ |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) |

Schedule 1

| T ₁ | T ₂ |
|--------------------------------------|---|
| read(A) $A := A - 50$ write(A) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) |
| read(B) $B := B + 50$ write(B) | read(B) $B := B + temp$ write(B) |

Schedule 3

Conflict Serializability

- Finally, consider **Schedule 7** of Figure:

| T_3 | T_4 |
|--------------|--------------|
| read(Q) | write(Q) |
| write(Q) | |

- It consists of only the significant operations (that is, the read and write) of transactions **T_3** and **T_4** .
- This schedule is **not conflict serializable**, since it is not equivalent to either the serial schedule $\langle T_3, T_4 \rangle$ or the serial schedule $\langle T_4, T_3 \rangle$.

Important Note

- It is possible to have **two schedules** that produce the **same outcome**, but that are **not conflict equivalent**.
- For **example**, consider **transaction T_5** , which transfers **\$10** from account B to account A.

Conflict Serializability

- Let **Schedule 8** be as defined in Figure.

| T_1 | T_5 |
|--------------------------------------|--------------------------------------|
| read(A) $A := A - 50$ write(A) | |
| | read(B) $B := B - 10$ write(B) |
| read(B) $B := B + 50$ write(B) | |
| | read(A) $A := A + 10$ write(A) |

Schedule 8

- We claim that **Schedule 8** is **not conflict equivalent** to the **serial schedule** , since, in **Schedule 8**, the **write(B)** instruction of **T5** **conflicts** with the **read(B)** instruction of **T1**.
- Thus, we cannot move all the instructions of **T1** before those of **T5** by swapping consecutive non-conflicting instructions.
- However, the final values of accounts **A** and **B** after the execution of either **schedule 8** or the **serial schedule** are the **same** —\$960 and \$2040, respectively.

Testing for Conflict Serializability


- Create a **precedence graph** $G(V,E)$ where V is the set of all the **transactions**.
- We will **create an edge** $T_i \rightarrow T_j$ if **one of the three condition is true**:

| T_i | T_j |
|---------------|---------------|
| W(Q) | |
| | R(Q) |

| T_i | T_j |
|---------------|---------------|
| R(Q) | |
| | W(Q) |

| T_i | T_j |
|---------------|---------------|
| W(Q) | |
| | W(Q) |

- The **schedule S** is **serializable** if and only if the **precedence graph** $G(V,E)$ has **no cycles**.



| T_3 | T_4 | T_6 |
|----------|----------|----------|
| read(Q) | write(Q) | |
| write(Q) | | write(Q) |

Schedule 9

Testing for Conflict Serializability

| T_1 | T_2 |
|--|---|
| read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y); | read_item(X); $X := X + M$; write_item(X); |

Schedule A

| T_1 | T_2 |
|--|---|
| read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y); | read_item(X); $X := X + M$; write_item(X); |





Schedule B

| T_1 | T_2 |
|--|---|
| read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y); | read_item(X); $X := X + M$; write_item(X); |

Schedule C


| T_1 | T_2 |
|--|---|
| read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y); | read_item(X); $X := X + M$; write_item(X); |

Schedule D



Testing for Conflict Serializability

| Transaction T_1 | Transaction T_2 | Transaction T_3 |
|--|--|--|
| <code>read_item(X);</code> <code>write_item(X);</code> <code>read_item(Y);</code> <code>write_item(Y);</code> | <code>read_item(Z);</code> <code>read_item(Y);</code> <code>write_item(Y);</code> <code>read_item(X);</code> <code>write_item(X);</code> | <code>read_item(Y);</code> <code>read_item(Z);</code> <code>write_item(Y);</code> <code>write_item(Z);</code> |
| Schedule E | | |



Equivalent serial schedules

None

Reason


Cycle $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

Cycle $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

Testing for Conflict Serializability

| Transaction T_1 | Transaction T_2 | Transaction T_3 |
|--|--|--|
| <code>read_item(X);</code> <code>write_item(X);</code> <code>read_item(Y);</code> <code>write_item(Y);</code> | <code>read_item(Z);</code> <code>read_item(Y);</code> <code>write_item(Y);</code> <code>read_item(X);</code> <code>write_item(X);</code> | <code>read_item(Y);</code> <code>read_item(Z);</code> <code>write_item(Y);</code> <code>write_item(Z);</code> |

Schedule F




Equivalent serial schedules

$T_3 \rightarrow T_1 \rightarrow T_2$

Serializability Order

- A **Serializability order** of the **Transactions** can be obtained through **Topological Sorting**.



View Serializability

- **View equivalence:**
- Consider two **schedules S** and **S'** , where the same set of **transactions** participates in both schedules.
- The **schedules S** and **S'** are said to be **view equivalent** if **three conditions** are met:
- **Condition 1:**
- For each **data item Q**, if transaction **T_i** reads the **initial value of Q** in **schedule S**, then **transaction T_i** must, in **schedule S'** , also **read the initial value of Q**.

| Schedule S | |
|----------------|----------------|
| T _i | T _j |
| Read(Q) | Read(Q) |
| | |

| Schedule S' | |
|----------------|----------------|
| T _i | T _j |
| Read(Q) | Read(Q) |
| | |

View Serializability

- **Condition 2:**
- For each **data item Q**, if transaction **T_i** executes **read(Q)** in **schedule S**, and if that value was produced by a **write(Q) operation** executed by **transaction T_j**, then the **read(Q) operation of transaction T_i** must, in **schedule S'**, also read the value of **Q** that was produced by the same **write(Q) operation** of **transaction T_j**.

| Schedule S | |
|----------------|----------------|
| T _i | T _j |
| Read(Q) | Write(Q) |

| Schedule S' | |
|----------------|----------------|
| T _i | T _j |
| Read(Q) | Write(Q) |

View Serializability

- **Condition 3:**
- For each **data item Q**, the transaction (if any) that performs the **final write(Q)** operation in **schedule S** must perform the **final write(Q)** operation in **schedule S'**.

| Schedule S | |
|---------------------|----------|
| Ti | Tj |
| Read(Q) Write(Q) | Write(Q) |
| Write(Q) | |

| Schedule S' | |
|---------------------|---------------------|
| Ti | Tj |
| Read(Q) Write(Q) | Read(Q) Write(Q) |
| | |

View Serializability

- **Conditions 1 and 2** ensure that each **transaction reads the same values** in both schedules and, therefore, performs the same computation.
- **Condition 3**, coupled with **conditions 1 and 2**, ensures that **both schedules result in the same final system state**.

| T ₁ | T ₂ |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) |

| T ₁ | T ₂ |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) |

- **Schedule 1 is not view equivalent to Schedule 2**, since, in **schedule 1**, the value of **account A** read by **transaction T2** was **produced by T1**, whereas this case does not hold in **Schedule 2**.

View Serializability

- However, **schedule 1** is **view equivalent** to **schedule 3**, because the values of account A and B read by transaction T2 were produced by T1 in both schedules.

| T ₁ | T ₂ |
|--|--|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> |

| T ₁ | T ₂ |
|--|--|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> |

- The concept of **view equivalence** leads to the concept of **View Serializability**.
- We say that a **schedule S** is **view serializable** if it is **view equivalent** to a **serial schedule**.

View Serializability

- **Schedule 9** is **view serializable** as it is **view equivalent** to the **serial schedule** $\langle T_3, T_4, T_6 \rangle$, since the one $\text{read}(Q)$ instruction reads the initial value of Q in both schedules, and T_6 performs the final write of Q in both schedules.

| T_3 | T_4 | T_6 |
|-------------------|-------------------|-------------------|
| $\text{read}(Q)$ | $\text{write}(Q)$ | |
| $\text{write}(Q)$ | | $\text{write}(Q)$ |

Schedule 9

- Every **Conflict-serializable** schedule is also **View serializable**, but there are **View serializable schedules** that are **not Conflict serializable**.
- Indeed, **Schedule 9** is **not conflict serializable**, since every pair of consecutive instructions conflicts, and, thus, **no swapping of instructions is possible**.
- In **Schedule 9**, transactions **T_4** and **T_6** perform **$\text{write}(Q)$ operations** without having performed a **$\text{read}(Q)$ operation**.
- Writes of this sort are called **blind writes**.
- **Blind writes** appear in any **view-serializable schedule** that is **not conflict serializable**.

Testing for View Serializability

- **Testing for View Serializability** is rather **complicated**.
- In fact, it has been shown that the **problem of testing for view serializability is itself NP-complete**.
- Thus, **almost certainly there exists no efficient algorithm to test for View Serializability**.

Recoverable Schedules

- Consider **schedule 11** in Figure below, in which **T9** is a transaction that performs only one instruction: **read(A)**.

| T ₈ | T ₉ |
|----------------|----------------|
| read(A) | |
| write(A) | |
| | read(A) |
| read(B) | |

- Suppose that the system allows **T9 to commit** immediately after executing the **read(A)** instruction.
- Thus, **T9 commits before T8** does.
- Now suppose that **T8 fails** before it commits.
- Since **T9 has read** the value of data item **A** written by **T8**, we must **abort T9** to ensure transaction **atomicity**.
- However, **T9 has already committed** and **cannot be aborted**.

Recoverable Schedules

- Thus, we have a situation where it is **impossible to recover correctly from the failure of T8.**
- **Schedule 11**, with the **commit** happening **immediately after the read(A)** instruction, is an **example** of a **Non-recoverable schedule** , which **should not be allowed**.
- A **Recoverable schedule** is one where, for **each pair of transactions Ti and Tj** such that :
 - **Tj reads a data item previously written by Ti.**
 - The **Commit operation of Ti** appears **before the Commit operation of Tj** .

| T ₈ | T ₉ |
|----------------|----------------|
| read(A) | |
| write(A) | |
| | read(A) |
| | write(A) |
| Commit | |
| | Commit |

Cascadeless Schedules

- Sometimes to **recover** correctly from the failure of a transaction T_i , we may have to **roll back several transactions**.
- **Example:** Consider the partial schedule of Figure:
- **Transaction T10** writes a value of **A** that is read by **transaction T11**.
- **Transaction T11** writes a value of **A** that is read by **transaction T12**.
- Suppose that, at this point, **T10 fails**.
- **T10 must be rolled back**.
- Since **T11 is dependent on T10**, **T11 must be rolled back**.
- Since **T12 is dependent on T11**, **T12 must be rolled back**.
- This **phenomenon**, in which a **single transaction failure leads to a series of transaction rollbacks**, is called **cascading rollback**.

| T_{10} | T_{11} | T_{12} |
|--------------------------------|---------------------|----------|
| read(A) read(B) write(A) | read(A) write(A) | read(A) |

Cascadeless Schedules

- **Cascading rollback is undesirable**, since it leads to the **undoing** of a significant amount of work.
- It is **desirable to restrict** the schedules to those where **cascading rollbacks** cannot occur.
- Such schedules are called **Cascadeless Schedules**.
- Formally, a **Cascadeless Schedule** is one where, for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- **Every Cascadeless schedule is also Recoverable.**