

Lecture 10

Concurrency Control

Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

Introduction

- One of the **fundamental properties** of a transaction is **Isolation**.
- When several **transactions** execute **concurrently** in the database, however, the **isolation property** may **no longer** be **preserved**.
- To ensure that it is, the system must **control the interaction** among the **concurrent transactions**.
- This **control** is achieved through one of a variety of **mechanisms** called **Concurrency-control schemes**.
- There are **two** main categories of **Concurrency-control schemes**:
 - **Lock-Based Protocols**
 - **Time-stamp Based Protocols**

Lock-Based Protocols

- One way to **ensure serializability** is to require that data items be accessed in a **mutually exclusive manner**.
- That is, while **one transaction is accessing a data item**, no other **transaction** can **modify** that data item.
- The most **common method** used to implement this requirement is to allow a transaction to access a **data item** only if it is currently **holding a lock on that item**.
- There are **two modes** in which a **data item** may be **locked**:
 - **Shared Mode**
 - **Exclusive Mode**

Lock-Based Protocols

1. Shared Mode

- If a **transaction T_i** has obtained a **shared-mode lock** (denoted by **S**) on item **Q**, then **T_i can read, but cannot write, Q.**

2. Exclusive Mode

- If a **transaction T_i** has obtained an **exclusive-mode lock** (denoted by **X**) on item **Q**, then **T_i can both read and write Q.**
- Every **transaction request a lock to Concurrency-control manager in an appropriate mode on data item Q**, depending on the types of **operations** that it will perform on Q.
- The **transaction can proceed** with the operation only after the **Concurrency-control manager grants the lock to the transaction.**

Lock-Based Protocols

- **Lock-compatibility matrix**
- The **compatibility** relation **between** the **two modes of locking** can be represented using a **matrix** .
- An **element comp(A, B)** of the **matrix** has the **value true** if and only if **mode A** is **compatible** with **mode B**.

	S	X
S	true	false
X	false	false
- Note that **shared mode** is **compatible** with **shared mode**, but **not** with **exclusive mode**.
- At any time, **several shared-mode locks** can be held **simultaneously** (by different transactions) on a **particular data item**.
- A subsequent **exclusive-mode lock** request has to **wait** until the **currently held shared-mode locks** are released.

Lock-Based Protocols

- A **transaction** requests a **shared lock** on **data item Q** by executing the **lock-S(Q)** instruction.
- Similarly, a **transaction requests an exclusive lock** through the **lock-X(Q)** instruction.
- A **transaction can unlock a data item Q** by the **unlock(Q)** instruction.
- **To access a data item, transaction T_i must first lock that item.**
- If the **data item is already locked by another transaction** in an **incompatible mode**, the **concurrency control manager will not grant the lock** until all **incompatible locks** held by other **transactions** have been **released**.
- **Thus, T_i is made to wait until all incompatible locks** held by other transactions **have been released.**

Example

T_1 : lock-X(B);
 read(B);
 $B := B - 50$;
 write(B);
 unlock(B);
 lock-X(A);
 read(A);
 $A := A + 50$;
 write(A);
 unlock(A).

T_2 : lock-S(A);
 read(A);
 unlock(A);
 lock-S(B);
 read(B);
 unlock(B);
 display(A + B).

T_1	T_2	concurrency-control manager
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T_2)
	unlock(A)	
	lock-S(B)	
	read(B)	grant-S(B, T_2)
	unlock(B)	
	display(A + B)	
lock-X(A)		
read(A)		
$A := A + 50$		
write(A)		
unlock(A)		grant-X(A, T_2)

- In this case, **transaction T2** displays **incorrect value**.
- The **reason** for this mistake is that the **transaction T1** unlocked data item B too **early**, as a result of which **T2** saw an **inconsistent state**.

Example

```
T3: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(B);  
    unlock(A).
```

```
T4: lock-S(A);  
    read(A);  
    lock-S(B);  
    read(B);  
    display(A + B);  
    unlock(A);  
    unlock(B).
```

- **Schedule corresponding to T3 and T4 will not lead to inconsistency.**

Lock-Based Protocols

- Unfortunately, **locking can lead to an undesirable situation.**
- Consider the **partial schedule** of Figure for **T3** and **T4**.

T ₃	T ₄
lock-X(B) read(B) B := B - 50 write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

- Since **T3** is holding an **exclusive-mode lock on B** and **T4** is requesting a **shared-mode lock on B**, **T4** is waiting for **T3** to unlock **B**.
- Similarly, since **T4** is holding a **shared-mode lock on A** and **T3** is requesting an **exclusive-mode lock on A**, **T3** is waiting for **T4** to unlock **A**.
- Thus, we have arrived at a **state where neither of these transactions can ever proceed with its normal execution.**
- This **situation** is called **Deadlock**.

Lock-Based Protocols

- When **deadlock occurs**, the **system must roll back** one of the two **transactions**.
- Once a **transaction has been rolled back**, the **data items** that were **locked by that transaction are unlocked**.
- These **data items** are **then available to the other transaction**, which can continue with its execution.
- Thus it require that **each transaction** in the system **follow a set of rules**, called a **locking protocol**, indicating when a transaction may lock and unlock each of the data items.

Starvation

- T2- lock-S(x) Granted
- T1- lock-X(x) Wait for T2
- T3- lock-S(x) Granted
- Commit T2 unlock(x)
- T1 still wait for T3 to unlock data item x.
- The transaction T1 may never make progress, and is said to be starved.
- We can avoid starvation of transactions by granting locks in the following manner:
- When a transaction T_i requests a lock on a data item Q in a particular mode M , the Concurrency-control manager grants the lock provided that :
 1. There is no other other transaction holding a lock on Q in a conflicting mode.
 2. There is no other transaction that is waiting for a lock on Q , and that made its lock request before T_i .

The Two-Phase Locking Protocol

- One **protocol** that **ensures serializability** is the **two-phase locking protocol**.
- This **protocol requires that each transaction issue lock and unlock requests in two phases:**
 - **1. Growing phase.** A transaction may **obtain locks**, but may **not release** any lock.
 - **2. Shrinking phase.** A transaction **may release locks**, but **may not obtain** any **new locks**.
- **Initially, a transaction is in the growing phase.**
- The **transaction acquires locks** as needed.
- **Once the transaction releases a lock, it enters the shrinking phase, and it can issue no more lock requests.**

The Two-Phase Locking Protocol

- Example: transactions T3 and T4 are two phase.

```
T3: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(B);  
    unlock(A).
```

```
T4: lock-S(A);  
    read(A);  
    lock-S(B);  
    read(B);  
    display(A + B);  
    unlock(A);  
    unlock(B).
```

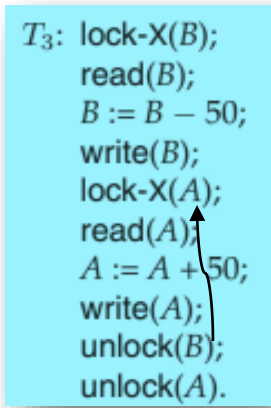
- On the other hand, transactions T1 and T2 are not two phase.

```
T1: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    unlock(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(A).
```

```
T2: lock-S(A);  
    read(A);  
    unlock(A);  
    lock-S(B);  
    read(B);  
    unlock(B);  
    display(A + B).
```

The Two-Phase Locking Protocol

- Note that the **unlock** instructions do not need to appear at the end of the transaction.
- **Example:** in the case of **transaction T3**, we could **move the unlock(B) instruction to just after the lock-X(A) instruction**, and **still retain the two-phase locking property**.



```
T3: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(B);  
    unlock(A).
```

- **Note:** the **two-phase locking protocol** ensures **Conflict serializability**.

The Two-Phase Locking Protocol

- **Cascading rollback** may occur under **two-phase locking**.
- Consider the **partial schedule** of Figure.

T_5	T_6	T_7
lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A)	lock-X(A) read(A) write(A) unlock(A)	lock-S(A) read(A)

- Each **transaction** observes the **two-phase locking protocol**, but the **failure of T5** after the **read(A)** step of **T7** leads to **cascading rollback** of **T6** and **T7**.

Strict two-phase locking protocol

- Cascading rollbacks can be avoided by a modification of two-phase locking called the **Strict two-phase locking protocol**.
- This **protocol** requires not only that locking be two phase, but also that **all exclusive-mode locks taken by a transaction be held until that transaction commits**.
- This requirement ensures that **any data written by an uncommitted transaction are locked in exclusive mode until the transaction commits**, preventing any other transaction from reading the data.
- **Another variant of two-phase locking is the Rigorous two-phase locking protocol**, which requires that **all locks be held until the transaction commits**.

Lock Conversion

```
T8: read(a1);  
    read(a2);  
    ...  
    read(an);  
    write(a1).
```

```
T9: read(a1);  
    read(a2);  
    display(a1 + a2).
```

- If we employ the **two-phase locking protocol**, then **T8 must lock a1 in exclusive mode**.
- Therefore, any **concurrent execution of both transactions** amounts to a **serial execution**.
- Notice, however, that **T8 needs an exclusive lock on a1 only at the end of its execution, when it writes a1**.
- Thus, if **T8 could initially lock a1 in shared mode**, and then **could later change the lock to exclusive mode**, we could get more concurrency, since **T8 and T9 could access a1 and a2 simultaneously**.

Lock Conversion

- This observation leads us to a **refinement of the basic two-phase locking protocol**, in which **lock conversions** are allowed.
- There is a **mechanism for upgrading a shared lock to an exclusive lock**, and **downgrading an exclusive lock to a shared lock**.
- Conversion from **shared to exclusive modes** denoted by **upgrade**, and from **exclusive to shared** by **downgrade**.
- **Lock conversion cannot be allowed arbitrarily.**
- Rather, **upgrading can take place** in only the **growing phase**, whereas **downgrading can take place** in only the **shrinking phase**.

T_8	T_9
lock-S(a_1)	lock-S(a_1)
lock-S(a_2)	lock-S(a_2)
lock-S(a_3)	
lock-S(a_4)	
	unlock(a_1)
	unlock(a_2)
lock-S(a_n)	
upgrade(a_1)	

Timestamp-Based Protocols for Concurrency Control

Timestamp-Based Protocols

- **Timestamps**
- With each **transaction T_i** in the system, we associate a **unique fixed timestamp**, denoted by **$TS(T_i)$** .
- This **timestamp** is **assigned by** the **database system** before the transaction **T_i** starts execution.
- If a **transaction T_i** has been assigned **timestamp $TS(T_i)$** , and a new **transaction T_j** enters the system, then **$TS(T_i) < TS(T_j)$** .
- There are **two** simple **methods** for implementing the **timestamp** scheme:
 - **System clock**
 - **Logical counter.**
- The **timestamps** of the **transactions** **determine** the **serializability order**.

Timestamp-Based Protocols

- Thus, if $TS(T_i) < TS(T_j)$, then the **system** must **ensure** that the produced **schedule** is **equivalent** to a **serial schedule** in which **transaction T_i** appears **before** **transaction T_j** .
- To **implement** this scheme, we associate with each **data item Q** **two timestamp values**:
 - **W-timestamp(Q)** denotes the **largest timestamp** of any **transaction** that **executed write(Q)** successfully.
 - **R-timestamp(Q)** denotes the **largest timestamp** of any **transaction** that **executed read(Q)** successfully.
- These **timestamps** are **updated** whenever a new **read(Q)** or **write(Q)** instruction is **executed**.

The Timestamp-Ordering Protocol

- The **Timestamp-ordering protocol** ensures that **any conflicting read and write operations** are **executed in timestamp order**.
- This **protocol** operates as follows:
- 1. Suppose that **transaction T_i** issues **$\text{read}(Q)$** .
 - a. If $\text{TS}(T_i) < \text{W-timestamp}(Q)$, then
 - **T_i needs to read a value of Q that was already overwritten.**
 - Hence, the **read operation is rejected**, and **T_i is rolled back**.
 - b. If $\text{TS}(T_i) \geq \text{W-timestamp}(Q)$, then
 - The **read operation is executed**, and
 - **$\text{R-timestamp}(Q) = \text{MAX}(\text{R-timestamp}(Q), \text{TS}(T_i))$.**

The Timestamp-Ordering Protocol

2. Suppose that **transaction T_i** issues **$\text{write}(Q)$** .

- a. If $\text{TS}(T_i) < \text{R-timestamp}(Q)$, then
 - the **value of Q** that **T_i** is **producing** was **needed previously**, and the **system assumed** that that **value would never be produced**.
 - Hence, the **system rejects** the **write operation** and **rolls T_i back**.
- b. If $\text{TS}(T_i) < \text{W-timestamp}(Q)$, then
 - **T_i is attempting to write an obsolete value of Q** .
 - Hence, the **system rejects** this **write operation** and **rolls T_i back**.
- c. Otherwise, the **system executes** the **write operation** and **sets $\text{W-timestamp}(Q)$ to $\text{TS}(T_i)$** .
- If a **transaction T_i** is **rolled back** by the **concurrency-control scheme** the **system assigns it a new timestamp** and **restarts it**.

The Timestamp-Ordering Protocol

T_{14} : read(B);
read(A);
display(A + B).

T_{15} : read(B);
 $B := B - 50$;
write(B);
read(A);
 $A := A + 50$;
write(A);
display(A + B).

T_{14}	T_{15}
read(B)	read(B)
	$B := B - 50$
	write(B)
read(A)	read(A)
display(A + B)	$A := A + 50$
	write(A)
	display(A + B)

- In **timestamp protocol**, a **transaction** is assigned a **timestamp** immediately **before** its **first instruction**.
- Thus, in **schedule** of Figure , **$TS(T_{14}) < TS(T_{15})$** .
- The **Timestamp-ordering protocol** ensures **Conflict Serializability**.
- This is because **conflicting operations** are **processed** in **timestamp order**.
- The **protocol ensures freedom from deadlock**, since **no transaction ever waits**.
- However, there is a **possibility** of **starvation** of long transactions if a **sequence** of **conflicting short transactions** causes **repeated restarting** of the long transaction.

Thomas' Write Rule

- **Thomas' Write Rule** is a **modification** to the **timestamp-ordering protocol** that allows **greater potential concurrency**.
- Let us consider **schedule** of Figure and **apply** the **timestamp-ordering protocol**.

T_{16}	T_{17}
read(Q)	write(Q)
write(Q)	

- Since **T16** starts before **T17**, we shall assume that $TS(T16) < TS(T17)$.
- The **read(Q)** operation of **T16** succeeds, as does the **write(Q)** operation of **T17**.
- When **T16** attempts its **write(Q)** operation, we find that $TS(T16) < W\text{-timestamp}(Q)$, since $W\text{-timestamp}(Q) = TS(T17)$.
- Thus, the **write(Q)** by **T16** is **rejected** and **transaction T16** must be **rolled back**.
- Any **transaction T_i** with $TS(T_i) < TS(T17)$ that **attempts a read(Q)** will be **rolled back**, since $TS(T_i) < W\text{-timestamp}(Q)$.

Thomas' Write Rule

- Any transaction T_j with $TS(T_j) > TS(T_{17})$ must read the value of Q written by T_{17} , rather than the value written by T_{16} .
- The **obsolete write operations can be ignored**.
- **Thomas' write rule:**
- Suppose that transaction T_i issues **write(Q)**.
 1. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was **previously needed**, and it had been assumed that the value would never be produced.
 - Hence, the **system rejects the write operation and rolls T_i back**.
 2. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is **attempting to write an obsolete value of Q** .
 - Hence, this **write operation can be ignored**.
 3. Otherwise, the system **executes the write operation** and sets $W\text{-timestamp}(Q)$ to $TS(T_i)$.

Deadlock Handling

- A **system** is in a **deadlock state** if there exists a **set of transactions** such that **every transaction** in the **set** is **waiting for another transaction** in the **set**.
- More precisely, there exists a **set of waiting transactions** $\{T_0, T_1, \dots, T_n\}$ such that
 - **T₀** is waiting for a data item that **T₁ holds**, and
 - **T₁** is waiting for a data item that **T₂ holds**, and ...,
 - .. and **T_{n-1}** is waiting for a data item that **T_n holds**, and
 - **T_n** is waiting for a data item that **T₀ holds**.
- **None of the transactions** can **make progress** in such a situation.
- The **only remedy** to this undesirable situation is for the system to **rolling back some of the transactions** involved in the **deadlock**.
- **Rollback** of a **transaction** may be **partial**: That is, a transaction may be rolled back to the point where it obtained a lock whose release resolves the **deadlock**.

Deadlock Handling

- There are **two principal methods** for dealing with the **deadlock problem**:
- **Deadlock Prevention**
 - Deadlock prevention protocol to ensure that the system will never enter a deadlock state.
- **Deadlock Detection and Recovery**
 - We can allow the system to enter a deadlock state, and
 - then try to recover by using a deadlock detection and deadlock recovery scheme.

Deadlock Prevention

- The simplest scheme of deadlock prevention requires that **each transaction locks all its data items** before it **begins execution**.
- Moreover, **either all are locked in one step or none are locked**.
- There are two main **disadvantages** to this protocol:
 - **(1)** it is often **hard to predict**, before the transaction begins, what data items need to be locked.
 - **(2) data-item utilization** may be **very low**, since many of the data items may be locked but unused for a long time.

Deadlock Prevention

- Two different **deadlock prevention schemes** using **timestamps** have been proposed:
- **1. Wait–die scheme (a non-preemptive technique)**
 - When **transaction T_i** requests a **data item** currently **held by T_j** ,
 - **T_i is allowed to wait only if $TS(T_i) < TS(T_j)$** (that is, T_i is older than T_j).
 - Otherwise, **T_i is rolled back** (dies).
- **Example**
- Suppose that transactions T_{22} , T_{23} , and T_{24} have timestamps 5, 10, and 15, respectively.
- If T_{22} requests a data item held by T_{23} , then **T_{22} will wait**.
- If T_{24} requests a data item held by T_{23} , then **T_{24} will be rolled back**.

Deadlock Prevention


- **Wound–wait scheme(preemptive technique)**
 - When **transaction T_i** requests a **data item** currently held by **T_j** ,
 - **T_i is allowed to wait** only if **$TS(T_i) > TS(T_j)$** (that is, **T_i is younger than T_j**).
 - **Otherwise, T_j is rolled back** (T_j is wounded by T_i).
- **Example**
- With transactions **T_{22} , T_{23} , and T_{24}** , if T_{22} requests a data item held by T_{23} , then the data item will be **preempted** from T_{23} , and T_{23} will be rolled back.
- If T_{24} requests a data item held by T_{23} , then T_{24} will wait.
- Whenever the system rolls back transactions, it is important to ensure that there is **no starvation**—that is, no transaction gets rolled back repeatedly and is never allowed to make progress.
- Both the **wound–wait** and the **wait–die** schemes **avoid starvation**.

Deadlock Detection

- **Deadlocks** can be described precisely in terms of a **directed graph** called a **wait-for graph**.
- This **graph** consists of a **pair $G = (V, E)$** , where V is a set of vertices and E is a set of edges.
- The **set of vertices** consists of **all the transactions** in the system.
- Each element in the set **E of edges** is an ordered pair **$T_i \rightarrow T_j$** .
- When **transaction T_i requests a data item currently being held by transaction T_j** , then the **edge $T_i \rightarrow T_j$** is inserted in the **wait-for graph**.
- **This edge is removed** only when transaction T_j is no longer holding a data item needed by transaction T_i .
- A **deadlock exists** in the system if and only if the **wait-for graph contains a cycle**.
- Each transaction involved in the cycle is said to be **deadlocked**.

Deadlock Detection


- **Example:**
- Consider the **wait-for graph** in Figure which depicts the following situation:



- Transaction T₂₅ is waiting for transactions T₂₆ and T₂₇.
- Transaction T₂₇ is waiting for transaction T₂₆.
- Transaction T₂₆ is waiting for transaction T₂₈.
- Since the **graph has no cycle**, the **system is not** in a **deadlock state**.

Deadlock Detection

- Suppose now that **transaction T28** is requesting an item held by **T27**.
- The **edge T28 → T27** is added to the **wait-for graph**, resulting in the new system state in Figure.



- This time, the **graph contains the cycle** $T_{26} \rightarrow T_{28} \rightarrow T_{27} \rightarrow T_{26}$ implying that **transactions T26, T27, and T28 are all deadlocked**.

Recovery from Deadlock

- The most common solution for Deadlock Recovery is to roll back one or more transactions to break the deadlock.
- Three actions need to be taken:
- **1. Selection of a victim.**
- Given a set of deadlocked transactions, we must **determine which transaction (or transactions) to roll back** to break the deadlock.
- We should **roll back those transactions** that will incur the **minimum cost**.
- **Many factors** may determine the **cost of a rollback**, including
 - **a.** How long the transaction has computed, and how much longer the transaction will compute before it completes its designated task.
 - **b.** How many data items the transaction has used.
 - **c.** How many more data items the transaction needs for it to complete.
 - **d.** How many transactions will be involved in the rollback.

Recovery from Deadlock

2. Rollback

- Once we have decided that a particular transaction must be rolled back, we must determine **how far this transaction should be rolled back**.
- The **simplest solution** is a **total rollback**: Abort the transaction and then restart it.
- However, it is **more effective to roll back the transaction only as far as necessary to break the deadlock**.

3. Starvation

- In a system where the selection of victims is based primarily on cost factors, it **may happen that the same transaction is always picked as a victim**.
- As a result, this **transaction never completes** its designated task, thus **there is starvation**.
- We must ensure that **transaction can be picked as a victim only a (small) finite number of times**.