

Lecture 4

Relational Model

Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

Introduction

- The **Relational Data Model** was first introduced by **Ted Codd** of **IBM Research** in **1970** in a classic paper (Codd 1970).
- **Relational model** attracted immediate attention due to its **simplicity** and **mathematical foundation**.
- The **model** uses the concept of a ***mathematical relation***—which looks somewhat like a **table** of values—as its **basic building block**.
- The **first commercial implementations** of the **Relational model** became available in the **early 1980s**, such as the **SQL/DS system by IBM** and the **Oracle DBMS**.
- Current **popular relational DBMSs** (RDBMSs) include **DB2** and **Informix Dynamic Server** (from IBM), **Oracle RDBMS**(from Oracle), **Sybase DBMS** (from Sybase) and **SQL Server** and **Access** (from Microsoft).
- In addition, several **open source systems**, such as **MySQL** and **PostgreSQL**, are available.

Relational Model Concepts

- The **Relational model** represents the **database** as a **collection of *relations***.
- Informally, each **relation** resembles a **table** of values.
- Each **row** in the **table** represents a **collection of related data values**.
- A **row** represents a fact that typically corresponds to a **real-world entity** or **relationship**.
- The **table name** and **column names** are used to help to interpret the meaning of the **values** in each **row**.
- For **example**, one **table** is named **STUDENT** because **each row** represents facts about a **particular student entity**.
- The **column names**—*Name*, *Student_number*, *Class*, and *Major*—specify how to interpret the **data values** in **each row**, based on the column each value is in.
- **All values in a column** are of the **same data type**.

Relational Model Concepts

STUDENT

Name	Student_number	Class	Major
Smith	17	1	CS
Brown	8	2	CS

COURSE

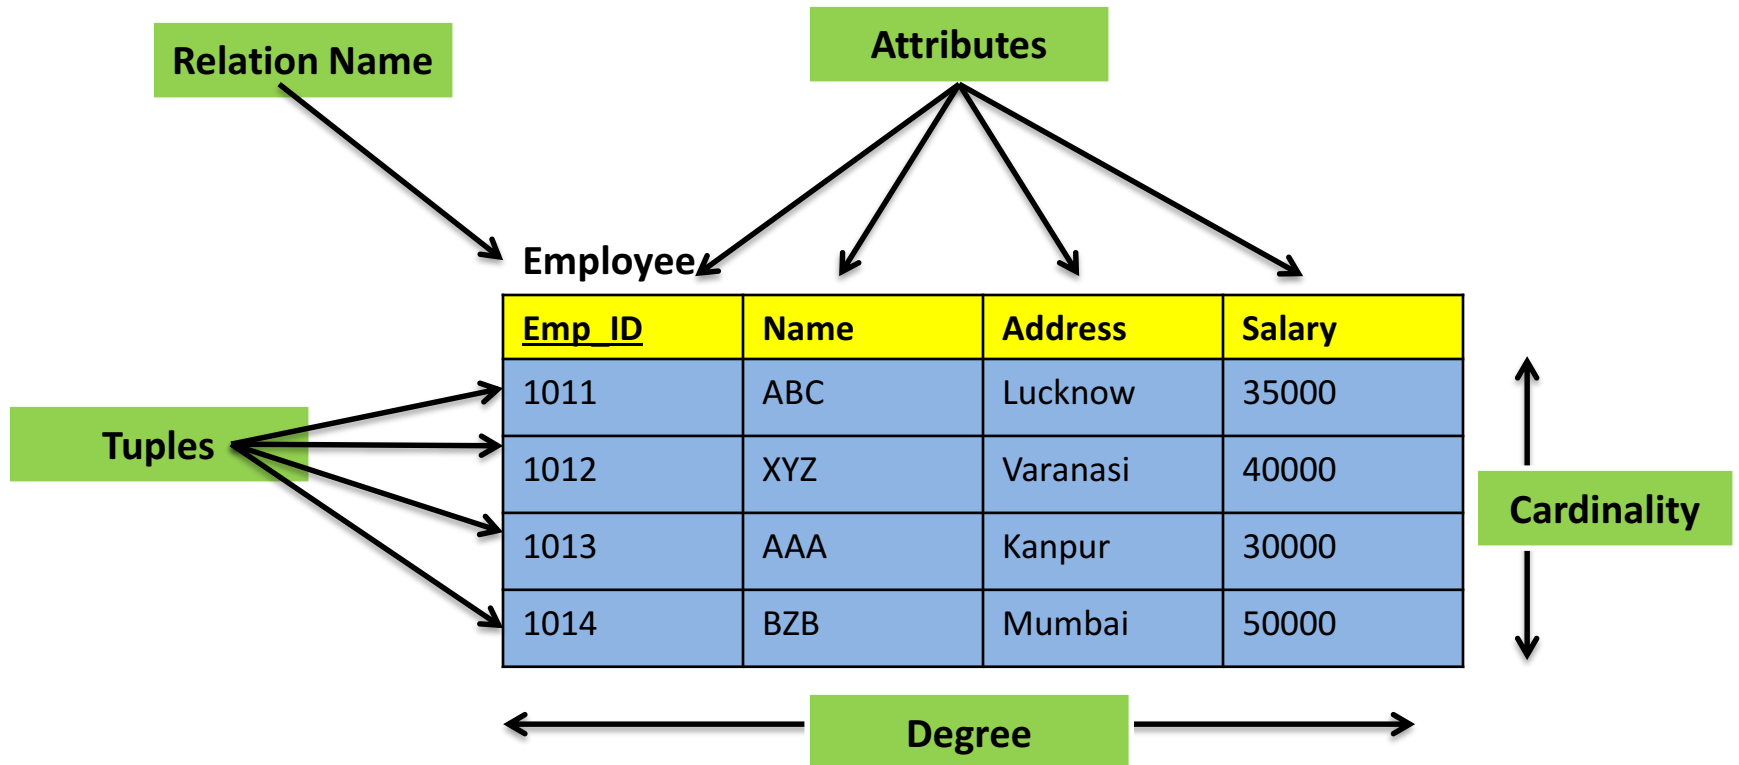
Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

Relation Schema

- A **relation schema** R , denoted by $R(A_1, A_2, \dots, A_n)$, is made up of a **relation name** R and a **list of attributes**, A_1, A_2, \dots, A_n .
- For each **attribute** A_i , there is a **set of permitted values**, called the **domain** D of that attribute in the relation schema R .
- D is called the **domain** of A_i and is denoted by $\text{dom}(A_i)$.
- A **relation schema** is used to **describe a relation**; R is called the **name** of this **relation**.
- The **Degree** of a **relation** is the **number of attributes** n of its relation schema.
- The **Cardinality** of a **relation** is the **number of tuples(records/rows)** in that relation.
- **Example:** A **relation** of **degree seven**, which stores information about university students, would contain **seven attributes** describing each student as follows:

***STUDENT**(Name, Ssn, Home_phone, Address, Office_phone, Age, Gpa)*

A Relation



Properties of a Relation

1. No duplicate tuples

- A **relation** can not contain two or more tuples which have the same values for all the **attributes**.
- Thus in any **relation every row(tuple)** should be **unique**.

2. Tuples are unordered

- The **order of rows(tuples)** in a **relation** is **immaterial**.

3. Attribute values are atomic

- Each **tuple** contains **exactly one value** for each **attribute**.

Relational Model Constraints

- There are various types of **Constraints** that can be specified on a relational database schema:
 1. **Domain Constraints**
 2. **Key Constraints**
 3. **Integrity Constraints**
 - a. **Entity Integrity Constraints**
 - b. **Referential Integrity Constraints**

1. Domain Constraints

- **Domain constraints** specify that within **each tuple**, the **value of each attribute A** in a **Relation** must be taken from the **same domain**.
- A **domain definition** usually consists of the following **components**:
 - *Data type*
 - *Size or length*
 - *Allowable values or allowable range*
- **Example:**
- **Names:** The set of character strings that represent names of persons.
- **Employee_ages:** Possible ages of employees in a company; each must be an integer value between 15 and 80.
- **Academic_department_names:** The set of academic department names in a university, such as Computer Science, Economics, and Physics.
- **Mobile_numbers:** The set of ten-digit numbers.

2. Key Constraints

- In the formal **Relational model**, a *relation* is defined as a *set of distinct tuples*.
- This means that **no two tuples** can have the **same combination of values** for **all their attributes**.
- The **subsets of attributes** of a **relation schema R** with the property that **no two tuples** in any **relation state r of R** should have the **same combination of values** for these **attributes**.
- Suppose that we denote **one such subset of attributes** by **SK**; then for **any two distinct tuples t_1 and t_2** in a **relation state r of R** , we have the **constraint** that:

$$t_1[\text{SK}] \neq t_2[\text{SK}]$$

- Any **such set of attributes SK** is called a **superkey** of the **relation schema R** .
- A **superkey SK** specifies a **uniqueness constraint** that no two distinct tuples in any **state r of R** can have the same value for **SK**.
- **Every relation** has **at least one default superkey**—the set of all its attributes.

2. Key Constraints

- A **key** K of a relation schema R is a **superkey** of R with the additional **property** that removing any attribute A from K leaves a set of attributes K that is **not a superkey** of R any more.
- Hence, a **key** satisfies **two properties**:
 1. **Two distinct tuples** in any state of the relation cannot have identical values for (all) the attributes in the key.
 - This first property also applies to a **superkey**.
 2. It is a **minimal superkey**—that is, a **superkey** from which we cannot remove any attributes and still have the uniqueness constraint in **condition 1** hold.
 - This property is not required by a **superkey**.
- A **key** is also a **superkey** but not **vice versa**.

2. Key Constraints

- Consider the **STUDENT** relation :
 - *STUDENT(Name, Ssn, Home_phone, Address, Age, Salary)*
- The **attribute set {Ssn}** is a **key** of **STUDENT** because no two student tuples can have the same value for Ssn.
- Any set of attributes that includes Ssn—for example, **{Ssn, Name, Age}**—is a **superkey**.
- However, the **superkey {Ssn, Name, Age}** is **not a key** of **STUDENT** because removing Name or Age or both from the set still leaves us with a **superkey**.
- In general, any **superkey** formed from a **single attribute** is also a **key**.
- A **key with multiple attributes** must require **all its attributes** together to have the **uniqueness property**.

2. Key Constraints

- In general, a **relation schema** may have **more than one key**.
- In this case, each of the **keys** is called a **candidate key**.
- For example, the **CAR relation** has **two candidate keys**: License_number and Engine_serial_number.
 - *CAR(License_number, Engine_serial_number, Make, Model, Year)*
- It is common to **designate one of the candidate keys** as the **primary key** of the **relation**.
- This is the **candidate key** whose values are used to *identify tuples* in the relation.
- We use the convention that the attributes that form the **primary key** of a **relation schema** are **underlined**.

Constraints on NULL Values

- Another **constraint on attributes** specifies whether **NULL values** are **permitted** or are **not permitted**.
- For **example**, if every **STUDENT tuple** must have a **valid, non-NULL** value for the **Name attribute**.
- Then **Name of STUDENT** is **constrained** to be **NOT NULL**.

Relational Database Schemas

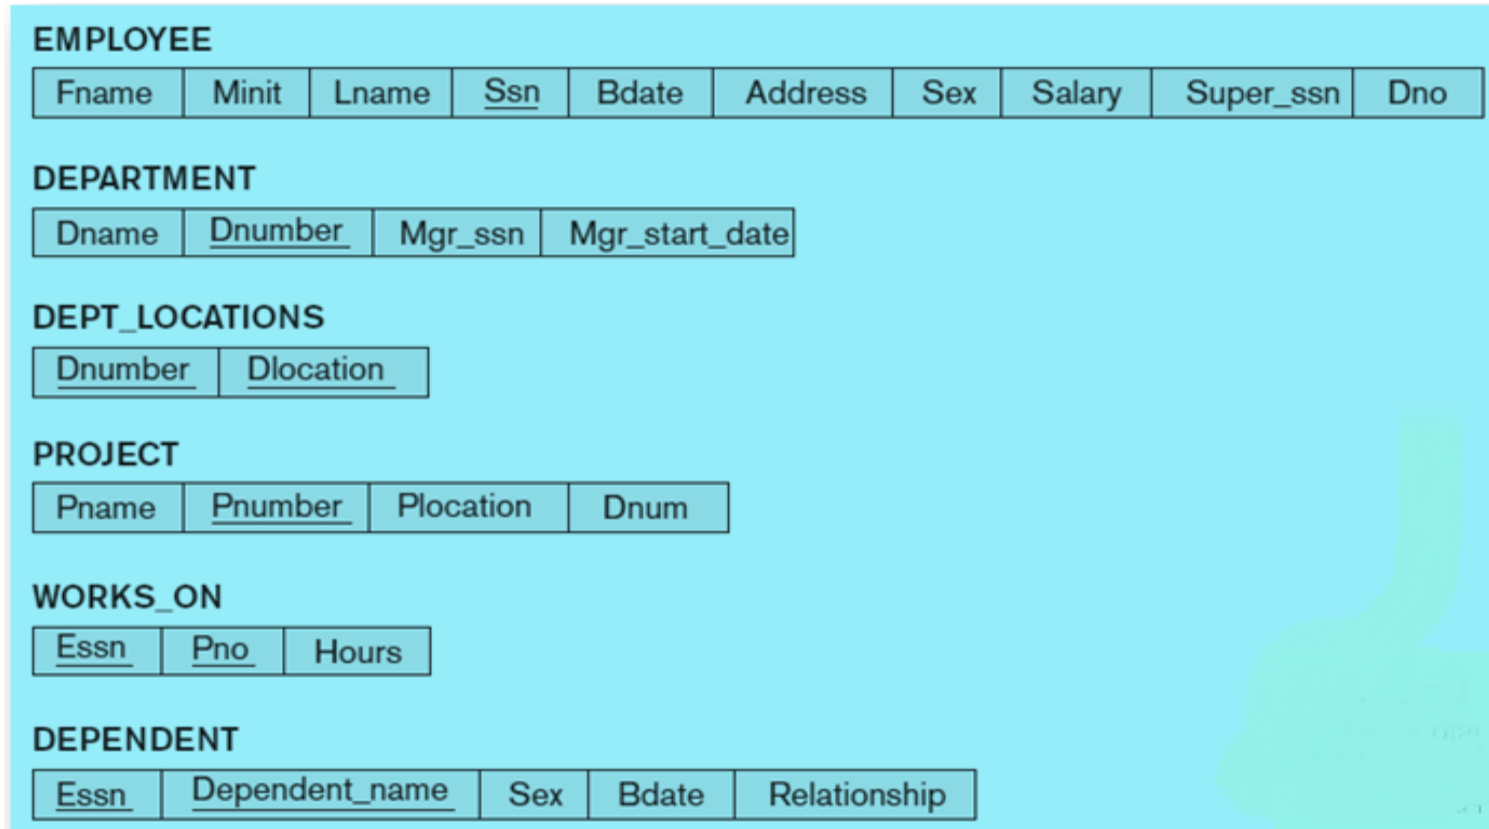
- A relational database schema S is a set of relation schemas $S = \{R_1, R_2, \dots, R_m\}$ and a set of integrity constraints IC .

- Example: A relational database schema

COMPANY = {EMPLOYEE, DEPARTMENT, DEPT_LOCATIONS, PROJECT, WORKS_ON, DEPENDENT}

- A relational database state DB of S is a set of relation states $DB = \{r_1, r_2, \dots, r_m\}$ such that each r_i is a state of R_i and such that the r_i relation states satisfy the integrity constraints specified in IC .
- A relational database implicitly include both its schema and its current state.
- A database state that does not obey all the integrity constraints is called an invalid state.
- A state that satisfies all the constraints in the defined set of integrity constraints IC is called a valid state.

Schema diagram



- Each **relational DBMS** must have a **Data definition language (DDL)** for **defining a relational database schema**.
- Current relational DBMSs are mostly using **SQL(Structured Query Language)** for this purpose.

Database state

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son

3. Integrity Constraints

- **Entity integrity constraint**
- The **entity integrity constraint** states that **no primary key value** can be **NULL**.
- This is because the **primary key value** is used to identify individual **tuples** in a **relation**.
- Having **NULL values** for the **primary key** implies that we cannot identify some tuples.
- For **example**, if two or more tuples had **NULL** for their **primary keys**, we may not be able to **distinguish** them if we try to reference them from other relations.
- **Key constraints** and **entity integrity** constraints are **specified on individual relations**.

3. Integrity Constraints

- **Referential integrity constraint**
- The **referential integrity constraint** is specified between **two relations** and is used to maintain the **consistency** among tuples in the **two relations**.
- The **referential integrity constraint** states that a **tuple** in one **relation** that **refers** to another relation **must refer** to an ***existing tuple*** in that **relation**.
- For **example**, the attribute **Dno** of **EMPLOYEE** gives the department number for which each employee works.
- Hence, its value in every **EMPLOYEE tuple** must match the **Dnumber value** of some tuple in the **DEPARTMENT relation**.
- **Foreign key** is used to specify the **referential integrity constraint** between the **two relation schemas R1 and R2**.

3. Integrity Constraints

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

Foreign key

- A set of attributes **FK** in relation schema **R1** is a foreign key of **R1** that references relation **R2** if it satisfies the following two rules:

Rule 1:

- The attributes in **FK** have the same domain(s) as the primary key attributes **PK** of **R2**; the attributes **FK** are said to reference or refer to the relation **R2**.

Rule 2:

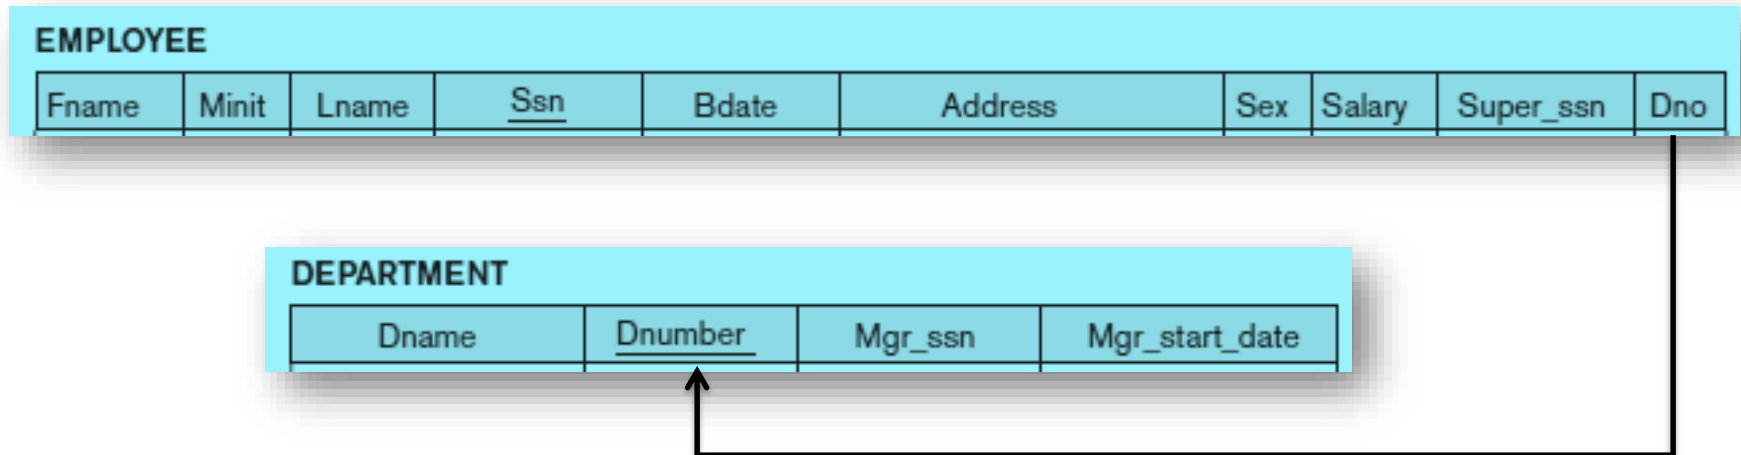
- A value of **FK** in a tuple **t1** of the current state **r1(R1)** either occurs as a value of **PK** for some tuple **t2** in the current state **r2(R2)** or is **NULL**.
 - In the former case, we have **t1[FK] = t2[PK]**, and we say that the tuple **t1** references or refers to the tuple **t2**.
- In this definition, **R1** is called the referencing relation and **R2** is the referenced relation.

Foreign key

- In a **Database of many relations**, there are usually **many referential integrity constraints**.
- **Referential integrity constraints** typically arise from the *relationships among the entities* .
- **Example:** In **COMPANY** database, the **EMPLOYEE** relation, the **attribute Dno** refers to the **department** for which an **employee works**; hence, we designate **Dno** to be a **foreign key** of **EMPLOYEE** referencing the **DEPARTMENT** relation.
- This means that a value of **Dno** in any **tuple t1** of the **EMPLOYEE** relation must **match** a value of the **primary key of DEPARTMENT**—the **Dnumber attribute**—in some **tuple t2** of the **DEPARTMENT** relation, or
- The value of **Dno** *can be NULL* if the employee does not belong to a department or will be assigned to a department later.

Foreign key

- We can *diagrammatically display referential integrity constraints* by drawing a directed arc from each foreign key to the relation it references.
- The **arrowhead** may **point to the primary key of the referenced relation**.



Self Referential Integrity

- A foreign key can *refer to its own relation*.
- Such **referential integrity constraint** is referred as **Self Referential Integrity constraint**.
- For **example**, the attribute **Super_ssn** in **EMPLOYEE** refers to the **supervisor** of an **employee**; this is another employee, represented by a tuple in the **EMPLOYEE** relation.
- Hence, **Super_ssn** is a **foreign key** that **references the EMPLOYEE relation itself**.
- In **EMPLOYEE** relation the tuple for employee 'John Smith' references the tuple for employee 'Franklin Wong,' indicating that 'Franklin Wong' is the supervisor of 'John Smith.'

Integrity Constraints

- All **integrity constraints** should be specified on the **relational database schema** (i.e., defined as part of its definition).
- Hence, the **DDL** includes **provisions for specifying** the various types of **constraints** so that the DBMS can automatically enforce them.
- Most **relational DBMSs** support **key, entity integrity**, and **referential integrity constraints**.
- These **constraints** are **specified** as a **part of data definition** in the **DDL**.

Referential integrity using FOREIGN KEY

- Referential integrity is specified via the **FOREIGN KEY** clause at **table level**, as shown below:

```
CREATE TABLE EMPLOYEE
```

```
(.....,
```

```
Ssn CHAR(9) NOT NULL,
```

```
Salary NUMBER(10,2),
```

```
Super_ssn CHAR(9),
```

```
Dno NUMBER NOT NULL,
```

```
PRIMARY KEY (Ssn),
```

```
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),
```

```
FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) );
```

Referential integrity using FOREIGN KEY

- A **Referential integrity constraint** can be **violated** when tuples are **inserted** or **deleted**, or when a foreign key or primary key attribute value is modified.
- The **default action** that SQL takes for an **integrity violation** is to **reject the update operation** that will cause a **violation**.
- However, the schema designer can specify an alternative action to be taken by attaching a **referential triggered action** clause to any **foreign key constraint**.
- The options include **SET NULL**, **CASCADE**, and **SET DEFAULT**.
- An option must be qualified with either **ON DELETE** or **ON UPDATE**.
- **Example:**
- The database designer chooses **ON DELETE SET NULL** and **ON UPDATE CASCADE** for the **foreign key Super_ssn** of **EMPLOYEE**.

Referential integrity using FOREIGN KEY

- What is a foreign key with **Cascade DELETE** in Oracle?
- A foreign key with **cascade delete** means that *if a record in the parent table is deleted, then the corresponding records in the child table will automatically be deleted.*
- A foreign key with a **cascade/set null** can be defined in either a **CREATE TABLE** statement or an **ALTER TABLE** statement.

- **Example:**

```
CREATE TABLE EMPLOYEE
```

```
(.....,
```

```
PRIMARY KEY (Ssn),
```

```
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn) ON DELETE SET NULL)
```

Referential integrity using FOREIGN KEY

- This means that if the tuple for a *supervisor employee* is *deleted*, the value of **Super_ssn** is automatically set to **NULL** for all employee tuples that were referencing the deleted employee tuple.
- On the other hand, if the **Ssn** value for a *supervisor employee* is *updated* (say, because it was entered incorrectly), the new value is *cascaded* to **Super_ssn** for all employee tuples referencing the updated employee tuple.

Example:

```
CREATE TABLE EMPLOYEE
```

```
( ...,
```

```
Dno INT NOT NULL DEFAULT 1, PRIMARY KEY (Ssn),
```

```
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn) ON DELETE SET NULL ON UPDATE  
CASCADE,
```

```
FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber) ON DELETE SET DEFAULT ON  
UPDATE CASCADE);
```

Referential integrity using FOREIGN KEY

ON DELETE	SET NULL	<i>Sets the rows in the child table to NULL when the corresponding row in the parent table is deleted.</i>
	SET DEFAULT	<i>Sets the rows in the child table to their default values if the corresponding rows in the parent table are deleted.</i>
	CASCADE	<i>Delete the rows from the child table automatically, when the rows from the parent table are deleted.</i>
ON UPDATE	SET NULL	<i>Sets the rows in the child table to NULL when the corresponding row in the parent table is updated.</i>
	SET DEFAULT	<i>Sets the rows in the child table to their default values if the corresponding rows in the parent table are updated.</i>
	CASCADE	<i>Update the rows from the child table automatically, when the rows from the parent table are updated.</i>